Model-Agnostic Meta Learning

Mid Eval Report

Sarthak Tangalpalli

241085

1   Introduction

Typically, what happens is that when we are machine learning models, we end up training them to accomplish only one given task by the use of loads of information. In a lab setting, this is not a problem, but when we get to the real world and then access information is a problem that is time-consuming and expensive. This is most evident in self-driving cars, wireless communication networks, as well as apps that try to recommend stuff to you.

$$\theta^* = \arg\min_\theta \mathbb{E}_{(x,y)\sim\mathcal{D}} \left[\mathcal{L}(f_\theta(x), y) \right]$$

The above expressions represent the standard supervised learning objective. The model f_θ is trained by minimizing the expected loss over a dataset $\mathcal{D}$. This formulation assumes a fixed task and a single optimal parameter vector θ^*, highlighting the limitation of task-specific learning and motivating the need for adaptive learning frameworks.

Humans function in a different way. We have an ability to adapt to new things very easily. We apply what we have to learn to learn new things in a very short span of time after observing very few examples. Hence, "learning to learn" emerged as a concept. The intention is to build an AI system that learns many things very efficiently and learns new unseen things in a very short span of time.

Meta-learning fills the gap. Rather than learning to excel at a single task, this approach concentrates its learning of general knowledge that can be applied across multiple tasks. Meta-learning is not solely concerned with seeking optimal parameters for solving one problem; instead, it envisions a dynamic learning foundation that can quickly adapt to any task. Meta-learning's concept lies at the core of model-agnostic meta-learning.

2   Tools and Experimental Setup

In this project, all programming was conducted using Python due to its inherent utility for scientific computing and AI as well as its being the language of choice for scientific computing and AI-related projects. The data handling tools that helped us deal with all the number-crunching that went on in this project are libraries such as NumPy and Pandas.

As far as constructing the models is concerned, we utilized software libraries capable of managing the calculations required for optimizing those models. This is a rather essential function when it comes to both learning models and meta-machine learning algorithms, upon which our project largely focused. Lastly, visualizing techniques were utilized to create graphics, enabling us to understand how models learn.

$$x_{\text{scaled}}^{(i)} = \frac{x^{(i)} - \mu}{\sigma}$$

$$\mu = \frac{1}{N} \sum_{i=1}^{N} x^{(i)}, \quad \sigma = \sqrt{\frac{1}{N} \sum_{i=1}^{N} (x^{(i)} - \mu)^2}$$

These equations describe feature standardization applied during data preprocessing. Each input feature is normalized using the dataset mean μ and standard deviation σ, ensuring numerical stability and consistent feature scaling. Such preprocessing improves convergence behavior for gradient-based optimization methods used in machine learning models.

Overall, this provided a very adaptable setup within which we were able to conduct our experiments. We were able to readily modify features associated with training and evaluation methods within our models to fully appreciate distinctions between traditional and meta-learning.

3   Supervised Learning Background

Supervised learning is a base for the majority of machine learning systems. In simple terms, it teaches a computer to connect inputs to the right answers using data that is already labeled. In this project, we looked at two main types: regression and classification.

Regression is used when we want to forecast exact numbers, such as signal strength. Simple "linear" regression attempts to find a straight-line relationship between the data. "Non-linear" regression can fit a greater number of patterns, including curves. However, it is more dangerous because it sometimes tries too hard and overfits every single point within the data-a type of problem called overfitting.

$$f_\theta(x) = \theta^T x$$

This equation defines a linear regression model, where predictions are computed as a linear combination of input features parameterized by θ.

Unlike the above-mentioned regression examples, Classification is different as it predicts categories, not numbers. For example, logistic regression calculates the probability that an object belongs to a specific group. Other techniques, such as decision trees and support vector machines, do classification by creating rules or boundaries separating data

into groups.

$$f_\theta(x) = \frac{1}{1 + e^{-\theta^T x}}$$

This equation defines the logistic regression model, where the sigmoid function maps linear outputs to probabilities in the range $(0, 1)$ for classification tasks.

All of these machine learning models learn by attempting a task a number of times and then testing how wrong they are by employing a strategy known as gradient descent to adjust their parameters. This is done repeatedly until these models learn to perform the task well.

$$\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}(f_\theta(x), y)$$

This expression describes the gradient descent update rule, where model parameters are iteratively updated in the direction that reduces the loss function.

4 Model Evaluation and Generalization

One of the biggest hurdles in Machine Learning is to ensure that the machine learning model generalizes well on unseen data that the model has not been trained on before. If the machine learning models are less complex, they are unable to identify significant patterns in the data (underfitting). However, if the models are more complex, there are chances that the models may memorize the data that is been trained on rather than actually learning (overfitting).

To figure out how well your model is performing, you need to keep track of various metrics. For regression problems where you predict a value, you can use tools such as Mean Absolute Error or Mean Squared Error. These will give you an idea of how accurate your predictions are and how strongly its outliers are impacting your predictions.

In classification tasks involving the grouping of objects for sorting, precision alone is no determinant of its quality. There is a more intricate assessment that involves metrics such as the confusion matrix, precision, and recall rate. All these provide an informed insight into the misclassification choices that the algorithm makes. The same conditions are applied to meta-learning as well.

$$\mathbb{E}_{\mathcal{D}}\left[(y - f_\theta(x))^2\right] = \text{Bias}^2 + \text{Variance} + \sigma^2$$

This expression decomposes the expected prediction error into bias, variance, and irreducible noise, explaining the tradeoff between model complexity and generalization.

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^{N} (y_i - f_\theta(x_i))^2$$

This equation defines the mean squared error, a commonly used metric for evaluating regression models by penalizing larger prediction errors more strongly.

5 Problem Formulation and Few-Shot Learning

Although standard artificial intelligence is perfectly effective if you have a huge, static dataset, it performs terribly if you don't have enough info or if your environment is dynamically changing. A great example is Wi-Fi networks, where obtaining precise signal information is costly and has to be done at a blazing pace, which is impossible for traditional artificial intelligence approaches.

It is also not very helpful to have to retrain the big model afresh every time the situation changes. It would be very time- and computationally-intensive. This is why we are seeing the end of "static" systems (learn and then do nothing else) and the beginning of "adaptive" systems (learn from the past and apply that knowledge when facing new situations).

But this is where Few-shot learning comes into play. Essentially, this is where one is asking a computer to learn a new task based solely upon a handful of examples. Meta-learning algorithms take this challenge head-on by creating models that can shift their focus with only a few tweaks.

$$\mathcal{T}_i = \left\{ \mathcal{D}_{\text{train}}^{(i)}, \mathcal{D}_{\text{test}}^{(i)} \right\}$$

This equation defines a task in the meta-learning framework as a pair of training and testing datasets.

$$\mathcal{D}_{\text{train}}^{(i)} = \{(x_j, y_j)\}_{j=1}^{K}$$

This expression represents a K-shot training dataset, emphasizing learning from a very small number of labeled examples.

6 Methodology: Model-Agnostic Meta-Learning

Model-Agnostic Meta-Learning (MAML) is a technique that has been developed with the aim of making AI models extremely flexible. The central approach might be described in the following way: If you begin in the right place, you just have to fine-tune the model slightly in order to be able to perform a brand new task.

This is how it all happens: Rather than practicing on individual examples of information, MAML practices on whole tasks. It selects a task to learn, attempting to learn that task from only a little bit of information available from that task. It then checks how well that task has been learned. Depending on that experience, it changes its starting position for its original task. It hopes to improve that starting position to the point that learning something new happens almost instantly.

$$\theta'_i = \theta - \alpha \nabla_\theta \mathcal{L}_{\mathcal{T}_i}(f_\theta)$$

This equation defines the inner-loop adaptation step of MAML, where task-specific parameters are obtained using gradient descent on the task loss.

$$\min_\theta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

This expression represents the meta-objective, which evaluates model performance after adaptation across multiple tasks.

$$\theta \leftarrow \theta - \beta \nabla_\theta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

This equation defines the meta-update step that optimizes the initial parameters to enable rapid adaptation on unseen tasks.

What MAML is best at is that it is a universal method. MAML does not care how your machine is made or how your problem is defined. Whether you are doing some number predictions or image sorting or training a robot to walk, MAML will do its job as it is not required to make tough assumptions.

7 Experiments and Implementation Details

The tasks involved in this project all centred on learning through action. We began with basics that included creating supervised models and evaluating results. Early experiments taught us why things like preventing overfitting and correctly rescaling your data actually make a difference in reality.

We then continued with meta-learning. This brought us the realization that there is a significant difference between training and "meta"-training. Through observing how the models learned for different tasks, we understood the importance of the "starting point" or initialization of parameters for how quickly a model can learn.

Despite the fact that a full-blown MAML environment is a resource-intensive process, these tasks have helped us gain insight into adaptation dynamics between models and the challenges that come into play when teaching a model to learn to transfer knowledge

from one task to another.

$$\hat{y}_i = f_\theta(x_i)$$

This expression computes the predicted output for an input sample using the current model parameters.

$$\mathcal{L}_{\text{train}} = \frac{1}{N} \sum_{i=1}^{N} (y_i - \hat{y}_i)^2$$

This equation defines the training loss minimized during experimental implementation.

8   Observations and Practical Applications

In all that we have explored, one thing that has become obvious is that adaptable models learn a lot faster with very fewer instances. Meta-learning shifts the training paradigm. The model is not only trained to complete a specific task; instead, it is trained to be super-sensitive to information. In this way, meta-learning enables the model to learn faster with only a few instances.

$$\|\theta'_i - \theta\| \ll \|\theta^* - \theta\|$$

This inequality indicates that only a small parameter update is required for task adaptation when starting from a meta-learned initialization.

$$\|\nabla_\theta \mathcal{L}_{\mathcal{T}_i}(f_\theta)\| \gg 0$$

This expression reflects the high sensitivity of the task loss to parameter changes, enabling fast convergence with minimal data.

This is incredibly valuable in real-world applications such as wireless networks or self-driving systems. This translates to less time spent collecting data, quicker acting times, and a system that is not brittle with changes in the environment. Though the process for meta-learning is more complex than traditional machine learning, its value in real-world applications is gigantic.

Analysis of MAML suggests that there has been a paradigm shift regarding the concept of AI in the sense that we are shifting our focus from experts that are only good for one task to systems that are broadly capable of dealing with whatever is next.

Conclusion

This report presents a detailed examination of the key concepts of machine learning and meta-learning, with a special focus on a concept called Model Agnostic Meta Learning (MAML).

Through the marriage of theory and experimentation, we were able to demonstrate exactly how one could go about designing these models that adapt on the fly. Ultimately, meta-learning proves to be a great tool in designing intelligent systems, particularly in those environments with limited amounts of data or those that are ever-changing.