

# Mid-Evaluation Report: Model Agnostic Meta Learning

**Daksh Battula**

Department of Electrical Engineering  
Indian Institute of Technology Kanpur

January 9, 2026

## Mid-Term Progress Report

*Foundational Python, Optimization in Machine Learning,  
and Model-Agnostic Meta-Learning Theory*

## Abstract

This document outlines the milestones achieved in the initial phase of the project, which focuses on designing a Few-Shot / Meta-Learning algorithm tailored for wireless communication systems. The core ambition is to facilitate rapid domain adaptation in dynamic network environments—such as millimeter-wave (mmWave) systems, Unmanned Aerial Vehicle (UAV) networks, and the Internet of Things (IoT)—where environmental conditions fluctuate frequently and large datasets are rarely available.

The report details the completion of essential technical modules, beginning with advanced Object-Oriented Programming (OOP) paradigms in Python and progressing to the manual implementation of foundational Machine Learning (ML) algorithms. A substantial section of this document investigates the theoretical mechanics of Stochastic Gradient Descent (SGD) and deep learning architectures. Furthermore, the report presents a comprehensive technical analysis of Model-Agnostic Meta-Learning (MAML), exploring its mathematical derivation, the "learning-to-learn" philosophy, and its specific utility in addressing channel estimation latency in 5G and 6G networks.

# 1 Introduction

The telecommunications sector is currently witnessing a major transition towards 5G and 6G standards. These next-generation infrastructures aim to deliver ultra-reliable, low-latency connectivity across diverse and complex environments. However, the physical constraints of these networks present substantial hurdles. High-frequency bands, such as mmWave, are notoriously susceptible to blockage and geometric interference. Similarly, a UAV navigating an urban canyon or an IoT sensor moved to a different room encounters a radio environment (channel state) that differs drastically from its training distribution.

Conventional Deep Learning (DL) methodologies struggle to cope with this variability. Standard models operate on the premise that training and testing data share an identical statistical distribution. When a wireless agent is deployed in a novel environment, traditional models demand retraining on massive new datasets. In wireless scenarios, gathering such data necessitates the transmission of thousands of "pilot signals," creating prohibitive bandwidth overhead and unacceptable delays.

## 1.1 Project Objective

The primary objective of this initiative is to negate the requirement for extensive re-training. We propose the application of **Few-Shot Learning**, specifically utilizing the Model-Agnostic Meta-Learning (MAML) framework. The aim is to train an agent that does not merely master a single task but instead learns an optimal *initialization* point. This allows the model to adapt to a novel environment using a minimal number of samples (e.g., fewer than 10 pilots).

This report summarizes the foundational work completed to date, covering structural software engineering (OOP), the mathematical engines of optimization (Gradient Descent), and the specific theoretical underpinnings of the MAML algorithm.

## 2 Module 1: Software Architecture and Object-Oriented Programming

Developing a robust Meta-Learning framework demands disciplined software engineering. The initial module focused on refreshing Object-Oriented Programming (OOP) concepts in Python to guarantee that the resulting codebase is modular, scalable, and easy to maintain.

### 2.1 Hierarchical Data Modeling

The assignment involved designing a hierarchical class structure to simulate a human population. While conceptually simple, this task was intended to reinforce state management and inheritance principles.

- **Base Class Implementation:** We established a parent class, `Human Being`, to bundle properties common to all entities, such as age and identity. This acts as the root of the inheritance tree.
- **Derived Classes:** We created `Teacher` and `Student` classes that inherit from the base. These introduced polymorphism—allowing methods with identical names to execute different logic based on the object type (e.g., a `work()` function might initiate grading for a teacher but study sessions for a student).
- **Nested Inheritance:** To demonstrate granular control, we further derived `Male` and `Female` subclasses within the specific roles, illustrating how attributes propagate through multiple levels of hierarchy.

### 2.2 Application to Meta-Learning

Although the assignment used demographic data, the architectural patterns are directly applicable to our project:

1. **Inheritance:** In our final system, we will design a generic `MetaLearner` parent class. Distinct algorithms (such as MAML or Reptile) will inherit from this base, utilizing shared utilities for logging and data handling while overriding specific `update_step` functions.
2. **Encapsulation:** Neural networks depend on precise state maintenance (weights and biases). OOP enables us to encapsulate these parameters, ensuring that "inner loop" adaptations do not inadvertently corrupt the "outer loop" meta-parameters during the learning process.

## 3 Module 2: Mathematical Foundations of Machine Learning

Prior to adopting high-level frameworks like PyTorch, it is essential to grasp the mathematical engines that drive learning. This module focused on implementing key algorithms from scratch using only NumPy.

### 3.1 Linear Regression and Feature Scaling

We built a Linear Regression model centered on the Mean Squared Error (MSE) cost function:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 \quad (1)$$

While analyzing a Real Estate dataset, we observed the issue of gradient instability caused by unscaled features. Variables with large numerical values (e.g., House Area) disproportionately influenced gradient updates compared to smaller features (e.g., Room Count). **Key Insight:** We applied Min-Max Normalization to compress all features into a  $[0, 1]$  range. This converted the error surface from a steep, elongated valley into a symmetrical bowl, enabling the Gradient Descent algorithm to converge with much greater speed and stability.

### 3.2 Logistic Regression and Probabilistic Classification

For the Breast Cancer classification task, we shifted to probabilistic modeling via Logistic Regression. The implementation hinged on the Sigmoid activation function:

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (2)$$

This exercise underscored the distinction in loss functions. Unlike regression, classification demands Log-Loss (Binary Cross-Entropy), which imposes heavy penalties on confident yet incorrect predictions. We manually calculated the partial derivatives for weights and biases, reinforcing the Chain Rule, which is the cornerstone of backpropagation.

### 3.3 K-Nearest Neighbors (KNN)

We investigated non-parametric learning using KNN on a Glass Identification dataset. Unlike previous models, KNN entails no training phase; all computation occurs during inference. **Critical Analysis:** We noted that KNN is extremely sensitive to the choice of  $k$ .

- **Low  $k$**  (e.g.,  $k = 1$ ): The model tends to overfit, reacting to noise and outliers.
- **High  $k$ :** The model underfits, blurring decision boundaries and losing the ability to distinguish local patterns.

## 4 Module 3: Advanced Optimization

### 4.1 Stochastic Gradient Descent (SGD) vs. Batch GD

A pivot component of this phase was understanding how models traverse the loss landscape. We contrasted Batch Gradient Descent (utilizing the entire dataset) with Stochastic Gradient Descent (SGD).

**Theoretical Finding:** Batch GD yields a smooth, deterministic path toward the minimum. However, in non-convex landscapes (typical in Deep Learning), it risks becoming trapped in local minima or saddle points. SGD, by updating parameters based

on individual samples, introduces "noise" into the optimization path.

$$\theta_{t+1} = \theta_t - \eta \nabla L(x_i, y_i; \theta_t) \quad (3)$$

This noise effectively acts as a regularizer, assisting the model in escaping shallow local minima to find more robust global solutions. This behavior is crucial for Meta-Learning, where the loss landscape is notoriously complex.

## 4.2 Decision Trees and Entropy

We examined Decision Trees through the lens of Information Theory. The fundamental goal of a tree is uncertainty reduction. We mathematically defined this uncertainty using **Entropy**:

$$H(S) = - \sum p_i \log_2 p_i \quad (4)$$

We determined that the optimal split is one that maximizes **Information Gain**—essentially selecting the attribute that divides the data into the most homogeneous child nodes possible.

## 4.3 Support Vector Machines (SVM)

We explored SVMs to grasp the concept of the "Maximum Margin" classifier. While Logistic Regression settles for *any* separating hyperplane, SVM seeks the *optimal* hyperplane that maximizes the gap to the nearest data points (support vectors).

**Mathematical Context:** To manage non-linear data (prevalent in complex wireless signals), we utilized the **Kernel Trick**. This method implicitly maps input vectors to a higher-dimensional space where they become linearly separable. A standard kernel we implemented is the Radial Basis Function (RBF):

$$K(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\sigma^2}\right) \quad (5)$$

### Key Learnings:

- **Sparsity:** Unlike models that depend on the full dataset, the SVM decision boundary is determined exclusively by the support vectors.
- **Generalization:** The max-margin principle inherently mitigates overfitting, rendering SVMs effective even when data is scarce.

## 5 Module 4: Deep Learning Frameworks (PyTorch)

Transitioning from manual NumPy implementations to PyTorch, we studied the mechanisms essential for deep neural networks.

- **Tensors and GPU Acceleration:** We mastered the manipulation of Tensors, which are generalized matrices optimized for GPU execution. This parallel processing is vital for the heavy matrix operations inherent in deep networks.

- **Autograd (Automatic Differentiation):** The most significant tool acquired was Autograd. For our upcoming Meta-Learning implementation, we must compute the derivative of a derivative (the Hessian). Performing this manually is error-prone. PyTorch’s dynamic computation graph automatically tracks gradients, enabling the implementation of sophisticated inner/outer loop optimization schemes.

## 6 Literature Review: Model-Agnostic Meta-Learning (MAML)

To establish a theoretical baseline for the project, we conducted a deep dive into the seminal paper on Model-Agnostic Meta-Learning (MAML). This section details the algorithmic mechanism we intend to adapt for the wireless domain.

### 6.1 The Problem: The ”Cold Start” in Wireless

In standard ML, a model  $f_\theta$  minimizes loss across a global dataset. However, in wireless communications, a ”task” corresponds to a specific physical environment. When a UAV enters a new area, the channel function  $H(x)$  shifts. A standard model fails because it assumes a static distribution. Retraining requires thousands of pilots, causing a ”Cold Start” lag where connection quality degrades significantly.

### 6.2 The MAML Solution: Learning to Initialize

MAML avoids learning a global model for all scenarios. Instead, it seeks a set of parameters  $\theta$  that are *highly adaptable*. The intuition is to locate a starting point on the loss landscape from which the optimal solution for *any* specific task is reachable within one or two gradient steps.

### 6.3 Algorithmic Derivation

MAML utilizes a bi-level optimization structure comprising an Inner Loop and an Outer Loop.

#### 6.3.1 The Inner Loop (Adaptation)

Assume we sample a task  $\mathcal{T}_i$  (e.g., a specific channel state). We are provided with  $K$  support samples (pilots). We calculate the gradients for this specific task and update the weights to generate temporary task-specific parameters  $\theta'_i$ :

$$\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta}) \quad (6)$$

Here,  $\alpha$  represents the inner learning rate. Crucially,  $\theta'_i$  is now a function of  $\theta$ .

#### 6.3.2 The Outer Loop (Meta-Optimization)

The objective of the meta-learner is to optimize the original  $\theta$  such that the *adapted* weights  $\theta'_i$  perform effectively on new query data (unseen data from the same environ-

ment). The meta-objective function is defined as:

$$\min_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \quad (7)$$

To minimize this, we execute a gradient descent step on  $\theta$ :

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \quad (8)$$

This is the pivotal step. We are computing the gradient of the loss (calculated at  $\theta'_i$ ) with respect to  $\theta$ . Since  $\theta'_i$  was derived using the gradient of  $\theta$ , this operation necessitates computing the gradient of a gradient—the Hessian matrix.

## 6.4 Computational Complexity and First-Order MAML

Full MAML requires second-order derivatives, which are computationally intensive ( $O(d^2)$ , where  $d$  is the parameter count). For wireless nodes with restricted processing capabilities, this presents a bottleneck. The literature proposes a "First-Order Approximation" (FO-MAML) where the second derivative term is omitted. Research suggests that FOMAML preserves much of the full MAML performance while significantly reducing computational burden, making it the ideal candidate for our IoT-based implementation.

## 7 Conclusion and Next Steps

This mid-term report verifies that the essential groundwork for the project is complete. We have progressed from fundamental software architecture to the implementation of sophisticated optimization algorithms and have successfully dissected the theoretical foundations of Meta-Learning.

The immediate next steps for the project include:

1. **Implementation:** Translating the mathematical formulation of MAML into a functional PyTorch module.
2. **Simulation:** Constructing a simulated wireless channel environment to generate "tasks" (varying channel realizations) for the meta-learner.
3. **Benchmarking:** Comparing the convergence rate of the MAML-initialized agent against a standard pre-trained neural network to quantify the improvement in adaptation speed (measured by the number of pilots required).