

Mid-Evaluation Report: Model Agnostic Meta Learning

Tejaswinee Rathore

Department of Electrical Engineering
Indian Institute of Technology Kanpur

January 9, 2026

Mid Eval Report

*Foundational Python, Machine Learning Optimization,
and Model-Agnostic Meta-Learning Theory*

Abstract

This report summarizes the progress made during the first phase of the project aimed at developing a Few-Shot / Meta-Learning-based training algorithm for wireless communications. The primary objective is to enable fast and efficient domain adaptation in modern network scenarios—specifically mmWave, Unmanned Aerial Vehicles (UAVs), and Internet of Things (IoT) systems—where environments change rapidly, and data scarcity is a persistent challenge.

This document details the completion of core technical modules, ranging from advanced Object-Oriented Programming (OOP) in Python to the implementation of fundamental Machine Learning (ML) algorithms from scratch. A significant portion of this report is dedicated to a theoretical analysis of Stochastic Gradient Descent (SGD) and deep learning frameworks. Finally, the report provides an in-depth technical review of Model-Agnostic Meta-Learning (MAML), dissecting its mathematical derivation, its "learning-to-learn" paradigm, and its specific potential to solve the channel estimation and latency problems inherent in 5G/6G networks.

1 Introduction

The telecommunications industry is currently undergoing a paradigm shift toward 5G and 6G technologies. These next-generation networks are expected to support hyper-reliable, low-latency communications across highly heterogeneous environments. However, the physical reality of these networks introduces significant complexity. High-frequency bands, such as millimeter-wave (mmWave), are extremely sensitive to blockage and environmental geometry. A UAV moving through a city or an IoT device entering a new room faces a drastically different radio environment (channel state) than what it was trained on.

Traditional Deep Learning (DL) approaches are ill-suited for this dynamic landscape. Standard models operate on the assumption that training and testing data are drawn from the same statistical distribution. When a wireless agent enters a new environment, a standard model requires retraining on a massive new dataset to adapt. In a wireless context, collecting this data requires transmitting thousands of "pilot signals." This overhead consumes valuable bandwidth and introduces unacceptable latency.

1.1 The Project Goal

The objective of this project is to eliminate the need for massive retraining. We propose utilizing **Few-Shot Learning**, specifically via the Model-Agnostic Meta-Learning (MAML) framework. The goal is to train an agent that does not learn a single task, but rather learns an *initialization* state that is capable of adapting to a new environment using only a handful of samples (e.g., fewer than 10 pilot signals).

This report outlines the foundational work completed to achieve this goal. It covers the structural software engineering principles (OOP), the mathematical backbone of optimization (Gradient Descent), and the specific theoretical mechanics of the MAML algorithm.

2 Module 1: Software Architecture and Object-Oriented Programming

To build a complex Meta-Learning system, rigorous software engineering practices are required. The first module focused on revisiting Object-Oriented Programming (OOP) in Python to ensure that the final codebase is modular, scalable, and maintainable.

2.1 Hierarchical Data Modeling

The assignment required the construction of a hierarchical class system representing a human population. This seemingly simple task was designed to enforce the principles of state management and inheritance.

- **Base Class Implementation:** We defined a parent class `Human Being` that encapsulates attributes universal to all entities, such as biological age and identification. This serves as the "root" of the inheritance tree.
- **Derived Classes:** We implemented `Teacher` and `Student` classes inheriting from the base. These classes introduced polymorphism—methods that share a name but behave differently depending on the object type (e.g., a `work()` method might trigger grading for a teacher but studying for a student).
- **Nested Inheritance:** Further granularity was achieved by deriving `Male` and `Female` subclasses within the roles, demonstrating how attributes can flow down multiple levels of a hierarchy.

2.2 Relevance to Meta-Learning

While the assignment utilized human demographics, the architectural patterns directly translate to our project structure:

1. **Inheritance:** In our final framework, we will define a generic `MetaLearner` parent class. Specific algorithms (like MAML or Reptile) will inherit from this base, sharing common utilities like data loading and logging while overriding the specific `update_step` methods.
2. **Encapsulation:** Neural networks rely on precise state management (weights and biases). OOP allows us to encapsulate these parameters, ensuring that the "inner loop" updates do not accidentally corrupt the "outer loop" meta-parameters during the adaptation phase.

3 Module 2: Mathematical Foundations of Machine Learning

Before utilizing high-level libraries like PyTorch, it is critical to understand the mathematical engines driving learning. This module involved implementing core algorithms from scratch using only NumPy.

3.1 Linear Regression and The Importance of Normalization

We implemented a Linear Regression model utilizing the Mean Squared Error (MSE) cost function:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 \quad (1)$$

During the analysis of a Real Estate dataset, we encountered the phenomenon of vanishing or exploding gradients due to unscaled features. Variables with large magnitudes (e.g., House Area in sq. ft) dominated the gradient updates compared to smaller features (e.g., Number of Rooms). **Key Insight:** We implemented Min-Max Normalization to scale all features into the $[0, 1]$ range. This transformed the error surface from an elongated valley (which is hard to navigate) into a spherical bowl, allowing the Gradient Descent algorithm to converge significantly faster.

3.2 Logistic Regression and Decision Boundaries

For the Breast Cancer classification task, we transitioned to probabilistic modeling using Logistic Regression. The core implementation relied on the Sigmoid activation function:

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (2)$$

This assignment highlighted the difference in loss functions. Unlike regression, classification requires the Log-Loss (Binary Cross-Entropy), which heavily penalizes confident but wrong predictions. We manually derived the derivatives for the weights and biases, reinforcing the concept of the Chain Rule which is central to backpropagation.

3.3 K-Nearest Neighbors (KNN)

We explored non-parametric learning via KNN on a Glass Identification dataset. Unlike the previous models, KNN has no training phase; computation happens entirely at inference time. **Critical Analysis:** We observed that KNN is highly sensitive to the hyperparameter k .

- **Low k** (e.g., $k = 1$): The model overfits, capturing noise and outliers in the data.
- **High k :** The model underfits, smoothing over the decision boundaries and losing local detail.

4 Module 3: Advanced Optimization

4.1 Stochastic Gradient Descent (SGD) vs. Batch GD

A critical part of this phase was understanding how models navigate the loss landscape. We compared Batch Gradient Descent (using the whole dataset) against Stochastic Gradient Descent (SGD).

Theoretical Finding: Batch GD produces a smooth, deterministic trajectory toward the minimum. However, in non-convex landscapes (common in Deep Learning), it is prone

to getting stuck in local minima (saddle points). SGD, by updating parameters based on a single sample, introduces "noise" into the trajectory.

$$\theta_{t+1} = \theta_t - \eta \nabla L(x_i, y_i; \theta_t) \quad (3)$$

This noise acts as a form of regularization, helping the model "jump" out of shallow local minima and find more robust solutions. This concept is vital for Meta-Learning, where the loss landscape is extremely complex.

4.2 Decision Trees and Entropy

We analyzed Decision Trees from an Information Theory perspective. The core mechanism of a tree is to reduce uncertainty. We mathematically quantified uncertainty using **Entropy**:

$$H(S) = - \sum p_i \log_2 p_i \quad (4)$$

We learned that the "best" split is the one that maximizes **Information Gain**—effectively choosing the question that separates the data into the purest possible child groups.

4.3 Support Vector Machines (SVM)

We studied SVMs to understand the principle of the "Maximum Margin" classifier. While Logistic Regression searches for *any* hyperplane that separates the classes, SVM searches for the *optimal* hyperplane that maximizes the distance to the nearest data points (the support vectors).

Mathematical Context: To handle non-linear data (which is common in complex wireless signals), we explored the **Kernel Trick**. This technique implicitly maps input vectors into a higher-dimensional feature space where they become linearly separable. A common kernel we implemented is the Radial Basis Function (RBF):

$$K(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\sigma^2}\right) \quad (5)$$

Key Learnings:

- **Sparsity:** Unlike other models that rely on all training data, the SVM decision boundary is defined solely by the support vectors.
- **Generalization:** The max-margin property inherently prevents overfitting, making SVMs robust even with smaller datasets.

5 Module 4: Deep Learning Frameworks (PyTorch)

Transitioning from scratch implementations to PyTorch, we explored the mechanisms required for deep neural networks.

- **Tensors and GPU Acceleration:** We learned to manipulate Tensors, which are generalized matrices capable of running on GPUs. This parallelization is necessary for the matrix multiplications in deep networks.

- **Autograd (Automatic Differentiation):** The most critical tool learned was Autograd. In our upcoming Meta-Learning implementation, we must calculate the derivative of a derivative (the Hessian). Doing this manually is prone to error. PyTorch’s dynamic computation graph allows us to track gradients automatically, enabling the implementation of complex inner/outer loop optimization strategies.

6 Literature Review: Model-Agnostic Meta-Learning (MAML)

To ground the project theoretically, a detailed analysis of the seminal paper on Model-Agnostic Meta-Learning (MAML) was conducted. This section explains the mechanism of the algorithm that we intend to implement for the wireless domain.

6.1 The Problem: The ”Cold Start” in Wireless

In a standard setting, a model f_θ is trained to minimize loss across a global dataset. However, in wireless communications, a ”task” is defined by a specific environment (e.g., a specific street canyon or indoor layout). When a UAV enters this new environment, the channel function $H(x)$ changes. A standard model fails here because it assumes a static distribution. Retraining from scratch requires thousands of pilots, which causes a ”Cold Start” lag where the link quality is poor.

6.2 The MAML Solution: Learning to Initialize

MAML does not try to learn a global model that works everywhere. Instead, it tries to find a set of parameters θ that are *easy to adapt*. The intuition is to find a location on the loss landscape from which the optimal solution for *any* specific task is just one or two gradient steps away.

6.3 Algorithmic Derivation

MAML involves a bi-level optimization process consisting of an Inner Loop and an Outer Loop.

6.3.1 The Inner Loop (Adaptation)

Assume we sample a task \mathcal{T}_i (e.g., a specific channel condition). We are given K support samples (pilots). We compute the gradients for this specific task and update the weights to create temporary task-specific weights θ'_i :

$$\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta}) \quad (6)$$

Here, α is the inner learning rate. Note that θ'_i is now dependent on θ .

6.3.2 The Outer Loop (Meta-Optimization)

The goal of the meta-learner is to optimize the original θ such that the *adapted* weights θ'_i perform well on new query data (data from the same environment not seen in the inner

loop). The meta-objective function is:

$$\min_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \quad (7)$$

To minimize this, we perform a gradient descent step on θ :

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \quad (8)$$

This is the critical step. We are taking the gradient of the loss calculated at θ'_i , with respect to θ . Because θ'_i was calculated using the gradient of θ , this step involves computing the gradient of a gradient—the Hessian matrix.

6.4 Computational Complexity and First-Order MAML

The full MAML algorithm requires second-order derivatives, which are computationally expensive ($O(d^2)$ where d is the number of parameters). For wireless nodes with limited processing power, this is a bottleneck. The literature suggests a "First-Order Approximation" (FOMAML) where the second derivative term is ignored. Research indicates that FOMAML retains much of the performance of full MAML while drastically reducing computational overhead, making it the prime candidate for our implementation on IoT devices.

7 Conclusion and Next Steps

This mid-evaluation report confirms that the necessary groundwork for the project has been laid. We have moved from basic software architecture to the implementation of complex optimization algorithms and successfully analyzed the theoretical underpinnings of Meta-Learning.

The immediate next steps for the project are:

1. **Implementation:** Translate the mathematical derivation of MAML into a PyTorch module.
2. **Simulation:** Develop a simulated wireless channel environment to generate "tasks" (different channel realizations) for the meta-learner.
3. **Benchmarking:** Compare the convergence speed of the MAML-initialized agent against a standard pre-trained neural network to quantify the gain in adaptation speed (measured in number of required pilots).